



SQL (Structured Query Language)

[Astrahan, Gray, Lindsay, Selinger, ...]

- IBM Sequel language developed as part of System R project at the IBM San Jose Research Laboratory
- Renamed Structured Query Language (SQL)
- ANSI and ISO standard SQL:
 - SQL-86
 - SQL-89
 - SQL-92
 - SQL:1999 (language name became Y2K compliant!)
 - SQL:2003
- Commercial systems offer most, if not all, SQL-92 features, plus varying feature sets from later standards and special proprietary features.
 - Not all examples here may work on your particular system.



© Nick Roussopoulos

33



Chapter 3: Introduction to SQL

- Overview of The SQL Query Language
- Data Definition
- Basic Query Structure
- Additional Basic Operations
- Set Operations
- Null Values
- Aggregate Functions
- Nested Subqueries
- Modification of the Database



© Nick Roussopoulos

34



Domain Types in SQL

- **char(*n*)** Fixed length character string, with user-specified length *n*.
- **varchar(*n*)** Variable length character strings, with maximum length *n*.
- **int.** Integer (a finite subset of the integers).
- **smallint** Small integer (a subset of the integer domain type).
- **numeric(*p,d*)** Fixed point number, with user-specified precision of *p* digits, with *d* digits to the right of decimal point.
- **real, double precision.** Floating point and double-precision floating point numbers.
- **float(*n*).** Floating point number, with user-specified precision of at least *n* digits.
- More domains are discussed in chapter 4 (date, currency, etc.)



35

© Nick Roussopoulos



Create Table Construct

- An SQL relation is defined using the **create table** command:

```
create table r (A1 D1, A2 D2, ..., An Dn,  
              (integrity-constraint1),  
              ...,  
              (integrity-constraintk))
```

- *r* is the name of the relation
- each *A_i* is an attribute name in the schema of relation *r*
- *D_i* is the data type of values in the domain of attribute *A_i*

Integrity constraints

- not null
- primary key (*A*₁, ..., *A_n*)
- foreign key (*A_m*, ..., *A_n*) references *r*



36

© Nick Roussopoulos



Integrity Constraints in Create Table

Example: Declare *ID* as the primary key for *instructor*

```
create table instructor (  
  ID          char(5),  
  name        varchar(20) not null,  
  dept_name   varchar(20),  
  salary      numeric(8,2),  
  primary key (ID),  
  foreign key (dept_name) references department)
```

primary key declaration ensures not null

```
insert into instructor values ('10211', 'Smith', 'Biology', 66000);  
insert into instructor values ('10211', null, 'Biology', 66000);
```



37

© Nick Roussopoulos



And a Few More Relation Definitions

- create table *student* (
 ID varchar(5) primary key,
 name varchar(20) not null,
 dept_name varchar(20),
 tot_cred numeric(3,0),
 foreign key (dept_name) references department);
 - create table *takes* (
 ID varchar(5),
 course_id varchar(8),
 sec_id varchar(8),
 semester varchar(6),
 year numeric(4,0),
 grade varchar(2),
 primary key (ID, course_id, sec_id, semester, year),
 foreign key (ID) references student,
 foreign key (course_id, sec_id, semester, year) references section);
- Can we remove sec_id from the primary key?



38

© Nick Roussopoulos



Drop and Alter Table Constructs

- **drop table**
- **alter table**
 - **alter table *r* add *A D***
 - where *A* is the name of the attribute to be added to relation *r* and *D* is the domain of *A*.
 - All tuples in the relation are assigned *null* as the value for the new attribute.
 - **alter table *r* drop *A***
 - where *A* is the name of an attribute of relation *r*
 - Dropping of attributes not supported by many databases.



39

© Nick Roussopoulos



Basic SQL Structure

- **Basic commands**
 - Retrieve: select
 - Update: insert, delete, update

Example

```
select  e.name  
from    emp e           [from emp as e]  
where   e.age >30;
```

Comments

- *e* is a tuple variable ranging over the emp relation
- a tuple variable followed by a dot and an attribute is an *indexed tuple variable* and specifies the corresponding attribute of the tuple
- what follows the select keyword is the *target list* (e.name)
- what follows the from is called the *tuple variable list* and consists of a list of relations and variable names
- what follows the where keyword is the *qualification clause*. It is an arbitrary Boolean expression. ⚠



40

© Nick Roussopoulos



SQL

- Basic format of select

```
select    [distinct] target_list
from      tuple_variable_list
where     qualification
[group by group_list_subset]
[order by target_list_subset];
```

- Semantics

- Form the Cartesian product of all relations in the from clause
- evaluate qualification: select the subset of the cartesian product that satisfies the qualification.
- apply grouping: partition the above subset into
- apply order and eliminate duplicates: if distinct, remove duplicate tuples and order the tuples according to the order by clause.
- evaluate target list: eliminate columns from above that are not in the target list



41

© Nick Roussopoulos



SQL

- We'll write several queries on the following relational schema in SQL.


```
sailors(sid,sname,rating)
boats(bid, bname,color)
reserve(sid,bid,date)
```

- (1) Find the names of sailors who have reserved boat #2

```
select  s.sname
from    sailors s, reserve r
where   s.sid=r.sid and r.bid=2
```

- (2) Find the names of sailors who have reserved red boat.

```
select s.sname
from   sailors s, boats b, reserve r
where  s.sid=r.sid and r.bid=b.bid and b.color="red"
```



42

© Nick Roussopoulos



SQL

sailors(sid,sname,rating)
boats(bid, bname,color)
reserve(sid,bid,date)

(3) find the colors of boats reserved by Pat

```
select  b.color
from    sailor s, boats b, reserve r
where   s.sid=r.sid and r.bid=b.bid and s.sname="Pat"
```

(4) Find the names of sailors who have reserved at least one boat.

```
select  s.sname
from    sailor s, reserve r
where   s.sid=r.sid
```



43

© Nick Roussopoulos



SQL

sailors(sid,sname,rating)
boats(bid, bname,color)
reserve(sid,bid,date)

(5) Find the names of sailor who have reserved a red or green boat

```
select  s.sname
from    sailors s, boats b, reserve r
where   s.sid=r.sid and r.bid=b.bid and (b.color="red" or b.color="green")
```

(6) Find the names of sailors who have reserved both a red and a green boat

```
select  s.sname
from    sailors s, boats b, reserve r, boats b2, reserve r2
where   s.sid=r.sid and r.bid=b.bid and b.color="red"
and     s.sid=r2.sid and r2.bid=b2.bid and b2.color="green"
```



44

© Nick Roussopoulos



SQL

sailors(sid,sname,rating)
boats(bid, bname,color)
reserve(sid,bid,date)

More on target lists:

- * is an abbreviation for all attributes in the from clause list
- each item in a target list can be as general as `attribute_name = expression`, where expression is any arithmetic or string expression over indexed tuple variables and constants. It can also contain some built-in function like `sqrt`, `sin`, `mod`, etc., as well as aggregates (coming up)

Example: With rating an integer from 1 to 10, this query "gives bonus" to persons who sailed two different boats on the same day

```
select  s.sname, rrating=s.rating+2
from    sailors s, reserve r, reserve r2
where   s.sid=r.sid and s.sid=r2.sid and r.date=r2.date and r.bid != r2.bid
```



45

© Nick Roussopoulos



SQL

sailors(sid,sname,rating)
boats(bid, bname,color)
reserve(sid,bid,date)

Qualifications: Each item in a qualification can be general as `expression=expression`.

Example:

```
select  name1=s1.sname, name2=s2.sname
from    sailors s1 s2
where   2*s1.rating=s2.rating-1
```



46

© Nick Roussopoulos



SQL

Further elaboration:

- Tuple variables can be implicit if the system can figure out which relation each attribute belongs to.
- Table names can be used as tuple variables.

Example: Find names, ages, and departments of employees who are over 40 and work on the first floor.

```
select  ename,age, emp.dname
from    emp, dept
where   age>40 and floor=1 and emp.dname=dept.dname
```



47

© Nick Roussopoulos



Natural Join

- Natural join matches tuples with the same values for all common attributes, and retains only one copy of each common column

```
select *
from instructor natural join teaches;
```

ID	name	dept_name	salary	course_id	sec_id	semester	year
10101	Srinivasan	Comp. Sci.	65000	CS-101	1	Fall	2009
10101	Srinivasan	Comp. Sci.	65000	CS-315	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	CS-347	1	Fall	2009
12121	Wu	Comp. Sci.	90000	FIN-201	1	Spring	2010
15151	Mozart	Music	40000	MU-199	1	Spring	2010
22222	Einstein	Physics	95000	PHY-101	1	Fall	2009
32343	Einstein	History	60000	HS-351	1	Spring	2010
45565	Katz	Comp. Sci.	75000	CS-101	1	Spring	2010
45565	Katz	Comp. Sci.	75000	CS-319	1	Spring	2010
76766	Crick	Biology	72000	BIO-101	1	Summer	2009
76766	Crick	Biology	72000	BIO-301	1	Summer	2010



48

© Nick Roussopoulos



Natural Join (Cont.)

- Danger in natural join in comparing unrelated attributes with same name
 - List the names of instructors along with the titles of courses that they teach
 - Incorrect version (equates `course.dept_name` with `instructor.dept_name`)

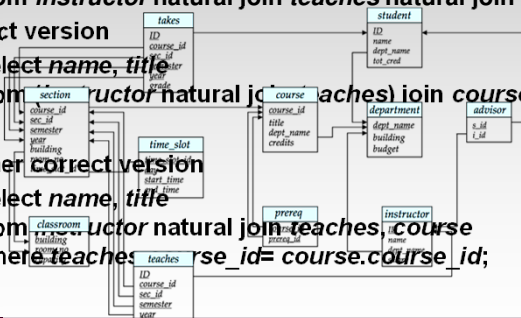
```
select name, title
from instructor natural join teaches natural join course;
```

- Correct version


```
select name, title
from instructor natural join teaches join course using(course_id)
```

- Another correct version


```
select name, title
from instructor natural join teaches course
where teaches.course_id = course.course_id;
```



49

© Nick Roussopoulos



The Rename Operation

- The SQL allows renaming relations and attributes using the `as` clause:
old-name as new-name

- E.g.,


```
select ID, name, salary/12 as monthly_salary
from instructor
```

- Find the names of all instructors who have a higher salary than some instructor in 'Comp. Sci'.


```
select distinct T.name
from instructor as T, instructor as S
where T.salary > S.salary and S.dept_name = 'Comp. Sci.'
```

- Keyword `as` is optional and may be omitted


```
instructor as T ≡ instructor T
```

50

© Nick Roussopoulos



String Operations

- SQL includes a string-matching operator for comparisons on character strings. The operator “like” uses patterns that are described using two special characters:
 - percent (%). The % character matches any substring.
 - underscore (_). The _ character matches any character.
- Find the names of all instructors whose name includes the substring “dar”.

```
select name
from instructor
where name like '%dar%'
```
- Match the string “100 %”

```
like '100 \%' escape '\'
```
- SQL supports a variety of string operations such as
 - concatenation (using “||”)
 - converting from upper to lower case (and vice versa)
 - finding string length, extracting substrings, etc.



51

© Nick Roussopoulos



Ordering the Display of Tuples

- List in alphabetic order the names of all instructors

```
select distinct name
from instructor
order by name
```
- We may specify **desc** for descending order or **asc** for ascending order, for each attribute; ascending order is the default.
 - Example: order by *name* desc
- Can sort on multiple attributes
 - Example: order by *dept_name*, *name*



52

© Nick Roussopoulos



Where Clause Predicates

- SQL includes a **between** comparison operator
- Example: Find the names of all instructors with salary between \$90,000 and \$100,000 (that is, $\geq \$90,000$ and $\leq \$100,000$)

```
select name
from instructor
where salary between 90000 and 100000;
```
- Tuple comparison

```
select name, course_id
from instructor, teaches
where (instructor.ID, dept_name) = (teaches.ID, 'Biology');
```



53

© Nick Roussopoulos



Duplicates

- In relations with duplicates, SQL treats the relations as multiset relations unless otherwise specified

```
select A1, A2, ..., An
from r1, r2, ..., rm
where P
```
- **Multiset** versions of some of the relational algebra operators – given multiset relations r_1 and r_2 :
 1. $\sigma_{\theta}(r_1)$: If there are c_1 copies of tuple t_1 in r_1 , that satisfy θ , then there are c_1 copies of t_1 in $\sigma_{\theta}(r_1)$.
 2. $\Pi_A(r)$: For each copy of tuple t_1 in r_1 , there is a copy of tuple $\Pi_A(t_1)$ in $\Pi_A(r_1)$ where $\Pi_A(t_1)$
 3. $r_1 \times r_2$: If there are c_1 copies of tuple t_1 in r_1 and c_2 copies of tuple t_2 in r_2 , there are $c_1 \times c_2$ copies of the tuple t_1, t_2 in $r_1 \times r_2$



54

© Nick Roussopoulos



Duplicates (Cont.)

- Example: Suppose multiset relations $r_1 (A, B)$ and $r_2 (C)$ are as follows:

$$r_1 = \{(1, a), (2, a)\} \quad r_2 = \{(2), (3), (3)\}$$

- $\Pi_B(r_1) = \{(a), (a)\}$
- $\Pi_B(r_1) \times r_2 = \{(a,2), (a,2), (a,3), (a,3), (a,3), (a,3)\}$



55

© Nick Roussopoulos



Null Values

- It is possible for tuples to have a null value, denoted by *null*, for some of their attributes
- *null* signifies an unknown value or that a value does not exist
- The result of any arithmetic expression involving *null* is *null*
 - Example: $5 + \text{null}$ returns null
- The predicate *is null* can be used to check for null values.
 - Example: Find all instructors whose salary is null.

```
select name
from instructor
where salary is null
```



56

© Nick Roussopoulos



Null Values and Three Valued Logic

- Any comparison with *null* returns *unknown*
 - Example: $5 < \text{null}$ or $\text{null} <= \text{null}$ or $\text{null} = \text{null}$
- Three-valued logic using the truth value *unknown*:
 - OR: $(\text{unknown or true}) = \text{true}$,
 $(\text{unknown or false}) = \text{unknown}$
 $(\text{unknown or unknown}) = \text{unknown}$
 - AND: $(\text{true and unknown}) = \text{unknown}$,
 $(\text{false and unknown}) = \text{false}$,
 $(\text{unknown and unknown}) = \text{unknown}$
 - NOT: $(\text{not unknown}) = \text{unknown}$
- Result of where clause predicate is treated as *false* if it evaluates to *unknown*



57

© Nick Roussopoulos



SQL Aggregates

- These functions operate on the multiset of values of a column of a relation, and return a value
 - avg: average value
 - min: minimum value
 - max: maximum value
 - sum: sum of values
 - count: number of values



58

© Nick Roussopoulos



SQL Aggregates

- Find the average salary of instructors in the Computer Science department
 - `select avg (salary)`
`from instructor`
`where dept_name= 'Comp. Sci.';`
- Find the total number of instructors who teach a course in the Spring 2010 semester
 - `select count (distinct ID)`
`from teaches`
`where semester = 'Spring' and year = 2010`
- Find the number of tuples in the *course* relation
 - `select count (*)`
`from course;`



59

© Nick Roussopoulos




Aggregate Functions – Group By

- Find the average salary of instructors in each department
 - `select dept_name, avg (salary) as avg_salary`
`from instructor`
`group by dept_name;`

ID	name	dept_name	salary
76766	Crick	Biology	72000
45565	Katz	Comp. Sci.	75000
10101	Srinivasan	Comp. Sci.	65000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000
12121	Wu	Finance	90000
76543	Singh	Finance	80000
32343	El Said	History	60000
58583	Califieri	History	62000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
22222	Einstein	Physics	95000

dept_name	avg_salary
Biology	72000
Comp. Sci.	77333
Elec. Eng.	80000
Finance	85000
History	61000
Music	40000
Physics	91000



60

© Nick Roussopoulos



Aggregation (Cont.)

- cannot mix tuple and aggregate values (attributes in select clause outside of aggregate functions must appear in the “group by” list)

```
select dept_name, ID, avg (salary)
from instructor
group by dept_name;
```

- multiple attribute “group by”
 - count the number of rooms used by course sections for each building per year

```
select    building, year, count(room_no)
from      section
group by  building, year
```



61

© Nick Roussopoulos



Aggregate Functions – Having Clause

- Find the names and average salaries of all departments whose average salary is greater than 42000

```
select dept_name, avg (salary)
from instructor
group by dept_name
having avg (salary) > 42000;
```

Note: predicates in the having clause are applied after the formation of groups whereas predicates in the where clause are applied before forming groups



62

© Nick Roussopoulos



Null Values and Aggregates

- Total all salaries

```
select sum (salary)
from instructor
```

 - Above statement ignores null amounts
 - Result is *null* if there is no non-null amount (all values are null)
- All aggregate operations except count(*) ignore tuples with null values on the aggregated attributes
- What if collection has only null values?
 - count returns 0
 - all other aggregates return null



63

© Nick Roussopoulos



Rownum

- select * from AUTHOR A
where rownum = 1;

ID	NAME
1	1 Anindya Datta
- select * from AUTHOR A
where rownum = 3;

ID	NAME
----	------
- select * from (select A.*, rownum as rn from AUTHOR A)
where rn = 3;

ID	NAME	RN
1	3 Sandeepan Banerjee	3
- select ID,Name from (select A.*, rownum as rn from AUTHOR A)
where rn = 3;

ID	NAME
1	3 Sandeepan Banerjee



64

© Nick Roussopoulos



SQL Set Operators

- union, intersect, and minus (except).

Example: Find the names of employees who work in the toy department and make at most 60K.

```
(select  ename
  from    emp
 where   dname="toy")
minus
(select  ename
  from    emp
 where   sal>60K)
```



65

© Nick Roussopoulos



SQL Nested Queries

Emp	ename	sal	dept	Dept	dname	floor	mgr
	Gary	30K	toy		candy	1	Irene
	Shirley	35K	candy		toy	2	Jim
	Christos	37K	shoe		men	2	John
	Robin	22K	toy		shoe	1	George
	Uma	30K	shoe				
	Tim	12K					

- A common use of subqueries is to perform tests for set membership, set comparisons, and set cardinality.
 - “in” tests for set membership
 - “not in” tests for non membership

Example: Find names of employees who work on the 1st floor.

```
select  ename
  from    emp
 where   dname in
           (select  dname
            from    dept
            where   floor=1)

select  ename
  from    emp
 where   dname not in
           (select  dname
            from    dept
            where   floor=1)
```

The same query in flat form is:

```
select  ename
  from    emp, dept
 where   emp.dname=dept.dname and floor=1
```



66

© Nick Roussopoulos



SQL

- Scoping of variables works exactly as in C

Example: Find the names who take a course from their advisor

```
select sname
from student
where s# in
      (select s#
       from enroll
       where c# in
            (select c#
             from class
             where prof=student.advisor))
```



67

© Nick Roussopoulos



SQL Connectives

- <op> any
- <op> all where <op> is one of { =, !=, <, >, >=, <= }

Example: Find names of employees who make more than everybody on the first floor.

```
select ename
from emp
where sal > all
      (select sal
       from emp, dept
       where emp.dname=dept.dname and floor=1)
```



68

© Nick Roussopoulos



Set Comparison

- Find names of instructors with salary greater than that of some (at least one) instructor in the Biology department.

```
select name
from instructor
where salary > some (select salary
                     from instructor
                     where dept name = 'Biology');
```

- Same query without “> some” clause

```
select distinct T.name
from instructor as T, instructor as S
where T.salary > S.salary and S.dept name = 'Biology';
```



69

© Nick Roussopoulos



Definition of Some Clause

- $F \langle \text{op} \rangle \text{some } r \Leftrightarrow \exists t \in r \text{ such that } (F \langle \text{op} \rangle t)$
Where $\langle \text{op} \rangle$ can be: $<, \leq, >, =, \neq$

$(5 < \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array}) = \text{true}$ (read: 5 < some tuple in the relation)

$(5 < \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{false}$

$(5 = \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{true}$

$(5 \neq \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{true (since } 0 \neq 5)$

Note: $(= \text{some}) \equiv \text{in}$ But $(\neq \text{some}) \neq \text{not in}$



70

© Nick Roussopoulos



Example Query

- Find the names of all instructors whose salary is greater than the salary of all instructors in the Biology department.

```
select name
from instructor
where salary > all (select salary
                    from instructor
                    where dept name = 'Biology');
```



71

© Nick Roussopoulos



Definition of all Clause

- $F <op> \text{all } r \Leftrightarrow \forall t \in r (F <op> t)$

(5 < all

0
5
6

) = false

(5 < all

6
10

) = true

(5 = all

4
5

) = false

(5 ≠ all

4
6

) = true (since 5 ≠ 4 and 5 ≠ 6)

Note (≠ all) is equivalent to not in see example 4 above

But (= all) is not equivalent to in see example 3 above



72

© Nick Roussopoulos



Test for Empty Relations

- The exists construct returns the value true if the argument subquery is nonempty.

```
select *  
from instructor  
exists (select *  
        from section  
        where year = '2020');
```

- not exists is true if subquery returns a nonempty result



73

© Nick Roussopoulos



Correlation Variables

- “Find all courses taught in both the Fall 2009 semester and in the Spring 2010 semester”

```
select course_id  
from section as S  
where semester = 'Fall' and year= 2009 and  
exists (select *  
        from section as T  
        where semester = 'Spring' and year= 2010  
        and S.course_id= T.course_id);
```

- Correlated subquery
- Correlation name or correlation variable



74

© Nick Roussopoulos



Not Exists

- Find all students who have taken all courses offered in the Biology department.

```
select distinct S.ID, S.name
from student as S
where not exists ( (select course_id
                   from course
                   where dept_name = 'Biology')
except
(select T.course_id
 from takes as T
 where S.ID = T.ID));
```

- Note that $X - Y = \emptyset \Leftrightarrow X \subseteq Y$
- Note: Cannot write this query using = all and its variants



75

© Nick Roussopoulos



Test for Absence of Duplicate Tuples

- The **unique** construct tests whether a subquery has any duplicate tuples in its result.
- Find all courses that were offered at most once in 2009

```
select T.course_id
from course as T
where unique (select R.course_id
              from section as R
              where T.course_id = R.course_id
              and R.year = 2009);
```



76

© Nick Roussopoulos



Derived Relations in the “from clause”


- SQL allows a subquery expression to be used in the “from clause”
- Find the average instructors’ salaries of those departments where the average salary is greater than \$42,000.”

```
select dept_name, avg_salary
from (select dept_name, avg (salary) as avg_salary
      from instructor
      group by dept_name)
where avg_salary > 42000;
```

- Note that we do not use the having clause

- Another way to write above query

```
select dept_name, avg_salary
from (select dept_name, avg (salary)
      from instructor
      group by dept_name) as dept_avg (dept_name, avg_salary)
where avg_salary > 42000;
```



77

© Nick Roussopoulos



Views

- In some cases, it is not desirable for all users to see the entire table
- Students need to know an instructors name and department, but not the salary.

ID	name	dept_name	salary
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Guthrie	Drama	87000

```
select ID, name, dept_name
from instructor
```

- Any relation other than the “base relational schema” but is made visible to a user via a query (“virtual relation”) is called a **view**.



78

© Nick Roussopoulos



View Definition

- A view is defined using the create view statement which has the form

create view *v* **as** < query expression >

where <query expression> is any legal SQL expression. The view name is represented by *v*.

- Once a view is defined, the view name can be used to refer to the virtual relation that the view generates.
- View definition is not the same as creating a new relation by evaluating the query expression
 - Rather, a view definition causes the saving the definition of a query expression;
- **Query Modification**: when a query uses a view, the defining expression is substituted in its place.



79

© Nick Roussopoulos



Example Views

- A view of instructors without their salary

```
create view faculty as  
select ID, name, dept_name  
from instructor;
```

- Create a view of department salary totals

```
create view departments_total_salary(dept_name, total_salary) as  
select dept_name, sum (salary)  
from instructor  
group by dept_name;
```



80

© Nick Roussopoulos



Views Defined Using Other Views

- **create view** *physics_fall_2009* **as**
select *course.course_id, sec_id, building, room_number*
from *course, section*
where *course.course_id = section.course_id*
and *course.dept_name = 'Physics'*
and *section.semester = 'Fall'*
and *section.year = '2009'*;
- **create view** *physics_fall_2009_watson* **as**
select *course_id, room_number*
from *physics_fall_2009*
where *building = 'Watson'*;



81

© Nick Roussopoulos



With Clause

- The **with** clause provides a way of defining a temporary view whose definition is available only to the query in which the with clause occurs.
- Find all departments with the maximum budget

with *max_budget (value)* **as**
 (select *max(budget)*
 from *department)*
select *budget*
from *department, max_budget*
where *department.budget = max_budget.value;*



82

© Nick Roussopoulos



Complex Queries using With Clause

- Find all departments where the total salary is greater than the average of the total salary at all departments

```
with dept_total(dept_name, value) as
  (select dept_name, sum(salary)
   from instructor
   group by dept_name),
dept_total_avg(value) as
  (select avg(value)
   from dept_total)
select dept_name
from dept_total, dept_total_avg
where dept_total.value >= dept_total_avg.value;
```



83

© Nick Roussopoulos



SQL Updates

- Insert command format:

insert into relation_name values (value_list)

or

insert into relation_name select_statement

- Semantics of insert:

- Format 1: Add the tuple corresponding to value_list into relation_name.
- Format 2: Execute the select statement, and then add the resulting tuples into relation_name.

Example:

```
insert into student
values(1, "Carey", "CS", "Stonebraker")
```



84

© Nick Roussopoulos



SQL Updates

Example: Suppose we have a relation
register(s#, name, paid),

in which registered students are recorded. After the end of registration week, we execute

```
insert into student
  select r.s#, r.name
  from   register r
  where  r.paid="yes"
```



85

© Nick Roussopoulos



SQL Delete

- Delete command format:
delete relation_name where qualification
- Semantics of delete: Execute the corresponding select command i.e.,

```
select full_target_list
  from relation_name
  where qualification
and then remove the resulting tuples from relation_name.
```



86

© Nick Roussopoulos



SQL Delete

The following command expels CS majors who received less than 1.5 in a CS course:

```
delete student
where major="CS" and s# in
      (select s# from enroll, course
       where enroll.s#student.s# and grade < 1.5
       and   enroll.c#course.c# and dept="CS")
```



87

© Nick Roussopoulos



SQL Update

➤ Update format

```
update relation_name
set   target_list
where qualification
```

1. Semantics of update: it is equivalent to executing the following:

```
1. insert into del_temp
   select full_target_list
   from relation_name
   where qualification

2. insert into app_temp
   select ext_target_list
   from relation_name
   where qualification
```

The difference of ext_target_list from full_target_list is that the former is augmented with tuple_variable.attr_name for all attributes of the from relation(s) that don't appear in full_target_list

```
3. remove the tuples in del_temp from relation_name
4. add the tuples in app_temp into relation_name.
```



88

© Nick Roussopoulos



SQL Update

Example: Give a 10% grade raise in every CS major in CS564.

```
update enroll
set grade=1.1 * grade
where c#="CS564" and s# in
      (select s#
       from student
       where major="CS")
```

The equivalent

```
insert into del_temp
select s#,c#,grade
from enroll
where c#="CS564" and s# in
      (select s#
       from student
       where major="CS")
```

```
insert into app_temp
select s#,c#,grade=1.1 * grade
from enroll
where c#="CS564" and s# in
      (select s#
       from student
       where major="CS")
```

then do (enroll minus del_temp) union app_temp



89

© Nick Roussopoulos



Outer Joins

- An extension of the join operation that avoids loss of information.
- Computes the join and then adds tuples from one relation that does not match tuples in the other relation to the result of the join.
- Uses *null* values.



90

© Nick Roussopoulos



Left Outer Join

Course

course_id	title	dept_name	credits
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

prereq

course_id	prereq_id
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

■ **course natural left outer join prereq**

course_id	title	dept_name	credits	prereq_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	null



91

© Nick Roussopoulos



Right Outer Join

Course

course_id	title	dept_name	credits
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

prereq

course_id	prereq_id
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

■ **course natural right outer join prereq**

course_id	title	dept_name	credits	prereq_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-347	null	null	null	CS-101



92

© Nick Roussopoulos



Full Outer Join

Course

course_id	title	dept_name	credits
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

prereq

course_id	prereq_id
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

■ **course natural full outer join prereq**

course_id	title	dept_name	credits	prereq_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	null
CS-347	null	null	null	CS-101



93

© Nick Roussopoulos



Update of a View

- Add a new tuple to *faculty* view which we defined earlier
insert into *faculty* values ('30765', 'Green', 'Music');
This insertion causes an insertion into the *instructor* relation of the tuple:
('30765', 'Green', 'Music', null)
- Most SQL implementations allow updates only on simple views
 - The key is in the view definition and the value is specified; other attributes can be omitted and are set to null.
 - No expressions, aggregates, or distinct specification in the select clause
 - No join, group by or having clause views
 - The query does not have a group by or having clause.



94

© Nick Roussopoulos



Integrity Constraints

- **Numerical constraints**
 - keys
 - quantified constraints of relationships (1-1, 1-N)
- **Domain constraints**
 - type declarations
 - fixed length char strings
 - [small] integer
 - floating point, double precision
 - extensions
 - date
 - currency
 - **NULL values**
 - for values to be filled in later on
 - primary keys cannot have NULL value
 - some attributes can be specified as NOT NULL
- **Key constraints**
 - keys have unique values
 - primary key- a candidate key declared primary
 - unique key- a candidate key
 - foreign key- a set of attributes that are primary keys for other relations)



95

© Nick Roussopoulos



Referential Integrity

- a value in a tuple of a relation must appear in at least one tuple of another relation, e.g. every value of eno in WORKS-IN is in at least one tuple of EMP
- EMP(eno,ename,sal)
DEPT(dno,dname,floor)
WORKS-IN(eno,dno,hours)
eno: foreign key dno: foreign key
- **referential integrity**
$$\pi_{\text{eno}}(\text{WORKS-IN}) \subseteq \pi_{\text{eno}}(\text{EMP})$$
 - **Updates may violate ref. integrity constraints**
 - insertion: inserts in the referencing relation a value that is not in the referenced one
 - update: referencing relation (as insertions)
 - deletion: delete a tuple referenced from the referencing relation
 - update: referenced relation (as deletions)



96

© Nick Roussopoulos



Assertions

- **declare arbitrary expressions that have to be satisfied**
assertion = predicate expression
e.g. GPA > 2.8
sum(all_charges) < credit_line
- **SQL:**
create assertion balance on deposit: balance >= 0
assert address on insertion deposit::
exists (select * from customer
where customer.customer.name = deposit.customer.name)
- When assertion is introduced, the condition is validated and if not rejected.
After that it stays a a watchdog for violations- VERY expensive



97

© Nick Roussopoulos



Triggers

- trigger is a statement that is executed when an update is applied to the database
 - the trigger has 2 parts
 - condition under which the trigger is executed
 - the actions to be taken when executed
- e.g.
- ```
define trigger overdraft on update of account T
(if new T.balance < 0 then (insert into borrow values
(T.branch.name,T.account-number,T.customer-name, T.balance)
update deposit S
set S.balance = 0
where S.account-number = T.account-number))
```
- triggers make the system reactive
  - triggers are also called active rules and they are also VERY expensive



98

© Nick Roussopoulos



## Large-Object Types

- Large objects (photos, videos, CAD files, etc.) are stored as a *large object*:
  - blob: binary large object -- object is a large collection of uninterpreted binary data (whose interpretation is left to an application outside of the database system)
  - clob: character large object -- object is a large collection of character data
  - When a query returns a large object, a pointer is returned rather than the large object itself.



99

© Nick Roussopoulos



## Authorization Specification in SQL

- The grant/revoke statement is used to authorize/de-authorize
  - grant select on (id,ename) on emp to Jeff,Jim,Joe;
  - grant update (adr) on emp to public;
  - grant delete on emp to Joe;
  - revoke select on emp from Jeff;
  - revoke update (adr) on emp from public;
- The format is
  - grant <privilege list> on <relation name or view name> to <user list>
  - <user list> is: a user-id, public
- Granting a privilege on a view does not imply granting any privileges on the underlying relations.
- The grantor of the privilege must already hold the privilege on the specified item (or be the dba).



100

© Nick Roussopoulos