



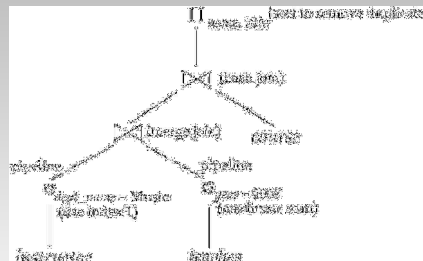
- **Cost-based optimization**
- **Dynamic Programming for Choosing Evaluation Plans**
- **Rule-Based Query Optimization**
- **Transformation of Relational Expressions**



© Nick Roussopoulos



- An evaluation plan defines exactly what algorithm is used for each operation, and how the execution of the operations is coordinated.



- **There are several- sometimes hundreds- of equivalent evaluation plans**
- **The optimizer evaluates the cost of all and selects the cheapest**
- **It relies on the ability to estimate both**
  - **the cost of each operation, and**
  - **the size of the results in each operation in order to evaluate the follow up operation**



© Nick Roussopoulos



## Selection $\sigma$ Size Estimation

➤ Case:  $\sigma_{A=v}(r)$

$$size = |\sigma_{A=v}| = \frac{n_r}{V(A, r)}$$

If **A** is a key, then  $V(A, r) = n_r$

$$size = |\sigma_{A=v}| = \frac{n_r}{n_r} = 1$$



© Nick Roussopoulos

322



## Selection $\sigma$ Size Estimation

➤ Case:  $\sigma_{A \leq v}(r)$  (case of  $\sigma_{A \geq v}(r)$  is symmetric)

- If  $\min(A, r)$  and  $\max(A, r)$  are available in catalog

$$size = |\sigma_{A \leq v}| = n_r * \frac{v - \min(A, r)}{\max(A, r) - \min(A, r)}$$

If  $v < \min(A, r)$  then  $size = |\sigma_{A \leq v}| = 0$

- If we know the value distribution of **A** (e.g. histograms) we can get much better estimates.
- In absence of statistical information *size* is assumed to be  $n_r/2$



© Nick Roussopoulos

323



## Size Estimation of Complex Selections

- The **selectivity** of a condition  $\theta_i$  is the probability that a tuple in the relation  $r$  satisfies  $\theta_i$
- If  $s_i$  is the number of tuples in  $r$  satisfying  $\theta_i$  the selectivity of it is  $s_i/n$
- **Conjunction result estimate** : (assuming Independence of predicates)

$$|\sigma_{\theta_1 \wedge \theta_2 \wedge \dots \wedge \theta_n}(r)| = n_r * \frac{s_1 * s_2 * \dots * s_n}{n_r^n}$$

- **Disjunction result estimate**:

$$|\sigma_{\theta_1 \vee \theta_2 \vee \dots \vee \theta_n}(r)| = n_r * \left( 1 - \left(1 - \frac{s_1}{n_r}\right) * \left(1 - \frac{s_2}{n_r}\right) * \dots * \left(1 - \frac{s_n}{n_r}\right) \right)$$

- **Negation**:  $\sigma_{\neg\theta}(r)$ .

$$|\sigma_{\neg\theta}(r)| = n_r - \text{size}(\sigma_{\theta}(r))$$



324

© Nick Roussopoulos



## Size Estimation of Join $n(R \bowtie S)$

$$0 \leq |n(r \bowtie s)| \leq n_r * n_s$$

- (0 is when nothing joins and  $n_r * n_s$  when everything joins)
- if joining attribute is a key of  $r$  then  $|n(r \bowtie s)| \leq n_s$   
each value of  $s.A$  would join to at most one value of  $r.A$
- if joining attribute is a key of  $r$  and a foreign key of  $s$  referencing  $r.A$  then  
 $|n(r \bowtie s)| = n_s$   
each value of  $s.A$  would join to exactly one value of  $r.A$
- if joining attribute is not a key then  
each value of  $A$  in  $r$  appears  $\frac{n_s}{V(A,s)}$  times in  $s$ , therefore,  
 $n_r$  tuples of  $r$  produce:  $|n(r \bowtie s)| = \frac{n_r * n_s}{V(A,s)}$

$$\text{symmetrically } n_s \text{ tuples of } s \text{ produce : } |n(r \bowtie s)| = \frac{n_s * n_r}{V(A,r)}$$

$$\text{if these two values are different we use: } \min\left\{\frac{n_r * n_s}{V(A,s)}, \frac{n_s * n_r}{V(A,r)}\right\}$$




325

© Nick Roussopoulos



## Join Example: student $\bowtie$ takes



Catalog stats:

$$n_{student} = 5,000 \quad f_{student} = 50 \quad b_{student} = 5,000/50 = 100$$

$$n_{takes} = 10,000 \quad f_{takes} = 25 \quad b_{takes} = 10,000/25 = 400$$

$V(ID, student) = 5,000$  (primary key)

$V(ID, takes) = 2,500$  (ID is a foreign key referencing student)

Because the join attribute ID is a FK

$$| student \bowtie takes | = n_{takes} = 10,000$$

$$\begin{aligned} \text{If ID was not a primary/foreign key in the above: } \min\left(\frac{n_t * n_s}{V(ID, student)}, \frac{n_s * n_t}{V(ID, takes)}\right) \\ = \min\left\{\frac{5,000 * 10,000}{5,000}, \frac{5,000 * 10,000}{2,500}\right\} = 10,000 \end{aligned}$$



326

© Nick Roussopoulos



## Size Estimation for Other Operations

- Projection: estimated size of  $\Pi_A(r) = V(A, r)$
- Aggregation : estimated size of  $\mathcal{A}_F(r) = V(A, r)$
- Set operations
  - For unions/intersections of selections on the same relation: use size estimate for selections
    - E.g.  $\sigma_{\theta_1}(r) \cup \sigma_{\theta_2}(r)$  can be rewritten as  $\sigma_{\theta_1 \vee \theta_2}(r)$
  - For operations on different relations:
    - estimated size of  $r \cup s$  = size of  $r$  + size of  $s$ .
    - estimated size of  $r \cap s$  = minimum size of  $r$  and size of  $s$ .
    - estimated size of  $r - s = r$ .

The above three estimates may be quite inaccurate, but provide upper bounds on the sizes.



327

© Nick Roussopoulos



## Size Estimation (Cont.)

### ➤ Outer joins:

- Estimated size of  $r \bowtie s = \text{size of } r \bowtie s + \text{size of } r$ 
  - Case of right outer join is symmetric  
 $= \text{size of } r \bowtie s + \text{size of } s$
- Estimated size of  $r \bowtie\!\!\!\! \sqsubset s = \text{size of } r \bowtie s + \text{size of } r + \text{size of } s$



328

© Nick Roussopoulos



## Join-Order Optimization

- Consider finding the best join-order for  $r_1 \bowtie r_2 \bowtie \dots \bowtie r_n$ .
- There are  $(2(n-1))/(n-1)!$  different join orders
  - With  $n = 7$ , the number is 665280,
  - With  $n = 10$ , the number is greater than 176 billion!
- No need to generate all the join orders
  - Using dynamic programming
  - the least-cost join order for any subset of  $\{r_1, r_2, \dots, r_n\}$  is computed only once and cached for its use in the algorithm (sub-optimal).




329

© Nick Roussopoulos



## Left Deep Join Trees

- In left-deep join trees, the right-hand-side input for each join is a base relation, not the result of an intermediate join
- Preferable because we have indices and better stats on base relations



330

© Nick Roussopoulos



## Choice of Evaluation Plans

- Must consider the interaction of evaluation techniques when choosing evaluation plans
  - choosing the cheapest algorithm for each operation independently may not yield best overall algorithm. E.g.
    - merge-join may be costlier than hash-join, but may provide a sorted output which reduces the cost for an outer level aggregation.
  - nested-loop join may provide opportunity for pipelining
- Practical query optimizers incorporate elements of the following two broad approaches:
  1. Search all the plans and choose the best plan in a cost-based fashion.
  2. Uses heuristics to choose a plan.



331

© Nick Roussopoulos



## Materialized Views

- A materialized view is a view whose contents are computed and stored
- Consider the view  

```
create view department_total_salary(dept_name, total_salary) as  
select dept_name, sum(salary)  
from instructor  
group by dept_name
```
- Materializing the above view would be very useful if the total salary by department is required frequently
  - Saves the effort of finding multiple tuples and adding up their amounts



332

© Nick Roussopoulos



## Materialized View Maintenance

- Keeping a materialized view up-to-date with the underlying data
- Materialized views can be maintained by re-computation on every update
- A much better option is to use incremental view maintenance
  - Changes to database relations are used to compute changes to the materialized view, which is then updated
- View maintenance can be done by
  - Manually defining triggers on insert, delete, and update of each relation in the view definition
  - Manually written code to update the view whenever database relations are updated
  - Periodic recomputation (e.g. nightly)
  - Lazy approach- incrementally update the view on demand (when accessed)
  - Some of the above methods are directly supported by many database system



333

© Nick Roussopoulos



## Incremental View Maintenance

- The changes (inserts and deletes) to a relation or expressions are referred to as its differential
  - Set of tuples inserted to and deleted from  $r$  are denoted  $i_r$  and  $d_r$
- To simplify our description, we only consider inserts and deletes
  - We replace updates to a tuple by deletion of the tuple followed by insertion of the update tuple
- We describe how to compute the change to the result of each relational operation, given changes to its inputs- define the **differential algebra**
- We then outline how to handle relational algebra expressions



334

© Nick Roussopoulos




## Join Operation


- Consider the materialized view  $v = r \bowtie s$



- Consider the case of an insert to  $r$ :



- Let  $r^{old}$  and  $r^{new}$  denote the old and new states of relation  $r$ 
  - We can write  $r^{new} \bowtie s$  as  $(r^{old} \cup i_r) \bowtie s$
  - And rewrite the above to  $(r^{old} \bowtie s) \cup (i_r \bowtie s)$
  - But  $(r^{old} \bowtie s)$  is simply the old value of the materialized view, so the incremental change to the view is just  $i_r \bowtie s$
- Thus, for inserts  $v^{new} = v^{old} \cup (i_r \bowtie s)$
- Similarly for deletes  $v^{new} = v^{old} - (d_r \bowtie s)$



335

© Nick Roussopoulos





## Selection and Projection Operations

- **Selection:** Consider a view  $v = \sigma_\theta(r)$ .

- $v^{new} = v^{old} \cup \sigma_\theta(I_r)$
- $v^{new} = v^{old} - \sigma_\theta(d_r)$

- **Projection is a more difficult operation**

- *Example:*  $R = (A, B)$ , and  $r(R) =$ 

a, 2
a, 3

$\Pi_A(r)$  has a single tuple (a).

If we delete the tuple (a,2) from  $r$ , we should not delete the tuple (a) from  $\Pi_A(r)$ , but if we then delete (a,3) as well, we should delete the tuple

- For each tuple in a projection  $\Pi_A(r)$ , keep a count of how many times it was derived
- On insert of a tuple to  $r$ , if the resultant tuple is already in  $\Pi_A(r)$  we increment its count, else we add a new tuple with count = 1
- On delete of a tuple from  $r$ , we decrement the count of the corresponding tuple in  $\Pi_A(r)$ 
  - if the count becomes 0, we delete the tuple from  $\Pi_A(r)$



336

© Nick Roussopoulos



## Aggregation Operations

- **count:**  $v = Ag_{count(B)}^{(r)}$ .

- When a set of tuples  $i_r$  is inserted, for each tuple  $r$  in  $i_r$ , if the corresponding group is already present in  $v$ , we increment its count, else we add a new tuple with count = 1
- When a set of tuples  $d_r$  is deleted, for each tuple  $t$  in  $i_r$ , we look for the group  $t.A$  in  $v$ , and subtract 1 from the count for the group.
  - When the count becomes 0, we delete the tuple from  $v$  for the group  $t.A$

- **sum:**  $v = Ag_{sum(B)}^{(r)}$

- We maintain the sum in a manner similar to count, except we add/subtract the  $B$  value instead of adding/subtracting 1 for the count
- Additionally we maintain the count in order to detect groups with no tuples. Such groups are deleted from  $v$ 
  - Cannot simply test for sum = 0 (why?)

- **sum:**  $v = Ag_{avg(B)}^{(r)}$

- we maintain the sum and count aggregate values separately, and divide at the end



337

© Nick Roussopoulos



## Aggregate Operations (Cont.)

- **Standard deviation:**  $v = \mathcal{A}g_{std(B)}(r)$ 
  - we maintain the sum of squares and the count
- **min, max:**  $v = \mathcal{A}g_{min(B)}(r)$ .
  - Handling insertions on  $r$  is straightforward.
  - Maintaining the aggregate values min and max on deletions may be more expensive. We have to look at the other tuples of  $r$  that are in the same group to find the new minimum
- **Percentiles and other non-distributive computations the problems are more difficult.**



338

© Nick Roussopoulos



## Other Operations

- **Set intersection:**  $v = r \cap s$ 
  - when a tuple is inserted in  $r$  we check if it is present in  $s$ , and if so we add it to  $v$ .
  - If the tuple is deleted from  $r$ , we delete it from the intersection if it is present.
  - Updates to  $s$  are symmetric
  - The other set operations, *union* and *set difference* are handled in a similar fashion.
- **Outer joins are handled in much the same way as joins but with some extra work (*bookkeeping*)**



339

© Nick Roussopoulos



## Query Optimization w/ Materialized Views

- **Rewriting queries to use materialized views:**
  - A user submits a query  $r \bowtie s \bowtie t$
  - A materialized view  $v = r \bowtie s$  is available
  - We can rewrite the query as  $v \bowtie t$
- **Alternatives**
  - Incremental update of the view and use
  - Recompute the view from its definition
  - Ignore the view and work with its base relations
- Query optimizer have been extended to consider all above alternatives and choose the best overall plan



340

© Nick Roussopoulos



## Materialized View Selection

- **Materialized view selection:** “What is the best set of views to materialize?”
- **Index selection:** “what is the best set of indices to create”
  - closely related, to materialized view selection but simpler
- **Materialized view selection and index selection based on typical system workload (queries and updates)**
  - Typical goal: minimize time to execute workload, subject to constraints on space and time taken for some critical queries/updates
  - One of the steps in database tuning
- Commercial database systems provide tools (called “tuning assistants” or “wizards”) to help the database administrator choose what indices and materialized views to create



341

© Nick Roussopoulos



## Top-K Query Optimization

- Queries may generate more than what is needed in a to-K query (i.e. all joinable tuples) but you only need 10

### Example

```
select * from r, s
where r.A = s.B
order by r.A ascending limit 10
```

- Alternative 1: use sort-merge join and stop at 10
- Alternative 2: make an estimate on the highest r.A value  $H$  in the 10 result tuples and modify the query to:  

```
select * from r, s
where r.A = s.B and r.A <= H
order by r.A ascending limit 10
```

  - If > 10 results, discard the extra (keep the 10)
  - If < 10 results, retry with larger  $H$



342

© Nick Roussopoulos



## Join Minimization

- Join minimization  $r(A,B), s(B,C,D)$   

```
select r.A, r.B
from r, s
where r.B = s.B
```
- Check if join with  $s$  is redundant, drop it
  - E.g. join condition is on foreign key from  $r$  to  $s$ ,  $r.B$  and no selection on  $s$ 
    - ```
select r.A, s2.B
from r, s s1, s s2
where r.B=s1.B and r.B = s2.B and s1.A < 20 and s2.A < 10
```
  - join with  $s1$  is redundant and can be dropped (along with selection on  $s1$ )
  - Lots of research in this area 70s/80s!



343

© Nick Roussopoulos



## Multiquery Optimization

➤ **Example**

**Q1: select \* from (r natural join t) natural join s**

**Q2: select \* from (r natural join u) natural join s**

- **Queries share common subexpression (r natural join s)**
- **May be useful to compute (r natural join s) once and use it in both queries (factor out)**

➤ **Multiquery optimization: find best overall plan for a set of queries, exploiting sharing of common subexpressions between queries where it is useful**

➤ **Implies the queries are synchronous**



344

© Nick Roussopoulos



## Materialized Views Optimization

- **Set of materialized views may share common subexpressions**
  - **view maintenance cost is shared**
- **The best approach is to maintain a Logical Access Path Schema (LAP schema) that captures all the relationships amongst the views**
- **Materialized view optimization/maintenance is amortized over a long period of time**



345

© Nick Roussopoulos



## Parametric Query Optimization

- **Example**  
select \*  
from r natural join s  
where r.a < \$1
  - value of parameter \$1 not known at compile time - only at run time
  - different plans may be optimal for different values of \$1
- **Solution 1: optimize at run time, each time query is submitted**
  - can be expensive
- **Solution 2: Parametric Query Optimization:**
  - optimizer generates a set of plans, optimal for different values of \$1
    - Set of optimal plans usually small for 1 to 3 parameters
    - Key issue: how to find set of optimal plans efficiently
  - best one from this set is chosen at run time when \$1 is known
- **Solution 3: Query Plan Caching**
  - If some plan is likely to be optimal for all parameter values, the optimizer caches the plan and reuses it, else reoptimizes each time
  - Implemented in many database systems




346

© Nick Roussopoulos



## 13.2 Rule-Based Optimization

- **Query:** Find the names of all instructors in the Music department, along with the titles of the courses that they taught in 2009.



- **Guideline: filter results as early as possible**
  - Work with most selective predicates first
  - Push selections ahead of joins
  - Order joins




347

© Nick Roussopoulos



## Equivalence of Relational Expressions

- Two expressions are equivalent if the result is the same (attributes and tuples)
  - even when tuples and attributes are ordered differently



- Which alternative is better?
  - The right one will generate less intermediate results. Is this always better?
    - for a single user environment and a single query, probably
    - In a multi-user multi-query optimization, maybe not



348

© Nick Roussopoulos



## Equivalence of Rel. Expressions and Rules

- Two relational algebra expressions are said to be *equivalent* if on every legal database instance the two expressions generate the same set of tuples
  - Again: order of tuples is irrelevant
- In SQL, inputs and outputs are multisets of tuples
  - Two expressions in the multiset version of the relational algebra are said to be equivalent if on every legal database instance the two expressions generate the same multiset of tuples
- An *equivalence rule* says that expressions of two forms are equivalent
  - Can replace expression of first form by second, or vice versa



349

© Nick Roussopoulos



## Equivalence Rules

1. Conjunctive selection operations can be deconstructed into a sequence of individual selections.

$$\sigma_{\theta_1 \wedge \theta_2}(E) = \sigma_{\theta_1}(\sigma_{\theta_2}(E))$$

2. Selection operations are commutative.

$$\sigma_{\theta_1}(\sigma_{\theta_2}(E)) = \sigma_{\theta_2}(\sigma_{\theta_1}(E))$$

3. Only the last in a sequence of projection operations is needed, the others can be omitted.

$$\Pi_{L_1}(\Pi_{L_2}(\dots(\Pi_{L_n}(E))\dots)) = \Pi_{L_1}(E)$$

4. Selections can be combined with Cartesian products and theta joins.

$$\text{a. } \sigma_{\theta}(E_1 \times E_2) = E_1 \bowtie_{\theta} E_2$$

$$\text{b. } \sigma_{\theta_1}(E_1 \bowtie_{\theta_2} E_2) = E_1 \bowtie_{\theta_1 \wedge \theta_2} E_2$$



350

© Nick Roussopoulos



## Equivalence Rules (Cont.)

5. Theta-join operations (and natural joins) are commutative:

$$E_1 \bowtie_{\theta} E_2 = E_2 \bowtie_{\theta} E_1$$

6. (a) Natural join operations are associative:

$$(E_1 \bowtie E_2) \bowtie E_3 = E_1 \bowtie (E_2 \bowtie E_3)$$

- (b) Theta joins are associative in the following manner:

$$(E_1 \bowtie_{\theta_1} E_2) \bowtie_{\theta_2 \wedge \theta_3} E_3 = E_1 \bowtie_{\theta_1 \wedge \theta_3} (E_2 \bowtie_{\theta_2} E_3)$$

where  $\theta_2$  involves attributes from only  $E_2$  and  $E_3$ .



351

© Nick Roussopoulos





## Equivalence Rules (Cont.)

7. The selection operation distributes over the theta join operation under the following two conditions:

(a) When all the attributes in  $\theta_0$  involve only the attributes of one of the expressions ( $E_1$ ) being joined.

$$\sigma_{\theta_0}(E_1 \bowtie_{\theta} E_2) = (\sigma_{\theta_0}(E_1)) \bowtie_{\theta} E_2$$

(b) When  $\theta_1$  involves only the attributes of  $E_1$  and  $\theta_2$  involves only the attributes of  $E_2$ .

$$\sigma_{\theta_1 \wedge \theta_2}(E_1 \bowtie_{\theta} E_2) = (\sigma_{\theta_1}(E_1)) \bowtie_{\theta} (\sigma_{\theta_2}(E_2))$$

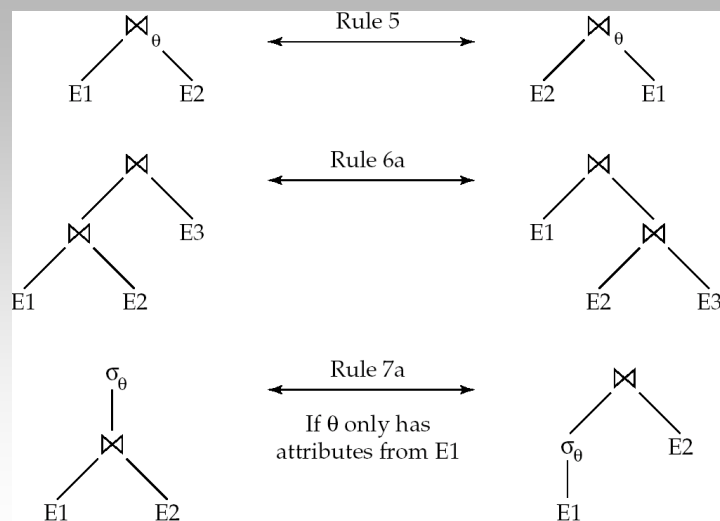


352

© Nick Roussopoulos



## Pictorial Depiction of Equivalence Rules



353

© Nick Roussopoulos



## Equivalence Rules (Cont.)

9. The set operations union and intersection are commutative

$$E_1 \cup E_2 = E_2 \cup E_1$$

$$E_1 \cap E_2 = E_2 \cap E_1$$

■ (set difference is not commutative).

10. Set union and intersection are associative.

$$(E_1 \cup E_2) \cup E_3 = E_1 \cup (E_2 \cup E_3)$$

$$(E_1 \cap E_2) \cap E_3 = E_1 \cap (E_2 \cap E_3)$$

11. The selection operation distributes over  $\cup$ ,  $\cap$  and  $-$ .

$$\sigma_{\theta}(E_1 - E_2) = \sigma_{\theta}(E_1) - \sigma_{\theta}(E_2)$$

and similarly for  $\cup$  and  $\cap$  in place of  $-$

Also: 
$$\sigma_{\theta}(E_1 - E_2) = \sigma_{\theta}(E_1) - E_2$$

and similarly for  $\cap$  in place of  $-$ , but not for  $\cup$

12. The projection operation distributes over union

$$\Pi_L(E_1 \cup E_2) = (\Pi_L(E_1)) \cup (\Pi_L(E_2))$$



354

© Nick Roussopoulos